Timings of common tasks using the **data.table** package in R

Matthew Dowle

Revised: November 7, 2011
(A later revision may be available on the homepage)

* WORK IN PROGRESS *

This document contains a series of tests, followed by a summary table of various timings and comparisons. Please go straight to the summary table first <here> in which each row has a link back to the test.

This document is reproducible. Simply run the .Rnw file yourself in your environment to confirm the results. Also see ?vignette, which says that edit(vignette("datatable-timings")) will extract the code from this document so you can easily work with it.

The .Rnw included in the package has N=10,000,000. This is a small number so that 'R CMD build' completes in a reasonable time (about 5 minutes). We don't want the nightly builds on R-Forge and CRAN to slow down just to run long timing comparisons. We have increased this to N=100,000,000 ourselves, and included the output on the datatable homepage (<link>).

Contents

1 Timing tests

1.1 Extraction

This is a repeat of the test in section 1 of the Introduction vignette. The syntax is explained there. This demonstrates the large difference in speed between vector scans and binary search. Therefore, please avoid using `==` in the `i` expression.

```
> n = ceiling(1e7/26^2)   # 10 million rows
> DF = data.frame(x=rep(LETTERS,each=26*n),
+                 y=rep(letters,each=n),
+                 v=rnorm(n*26^2))
> DT = data.table(DF,key="x,y")
> tables()

     NAME       NROW  MB COLS  KEY
[1,] DT   10,000,068 153 x,y,v x,y
Total: 153MB

> tt=system.time(ans1 <- DF[DF$x=="R" & DF$y=="h",]); tt
```

```
   user  system elapsed
 12.517   1.016  13.574

> head(ans1)

        x y           v
6642058 R h -0.8196355
6642059 R h  1.1127891
6642060 R h -0.2415036
6642061 R h  1.1443273
6642062 R h  1.4476976
6642063 R h  1.1084013

> dim(ans1)

[1] 14793      3

> ss=system.time(ans2 <- DT[J("R","h")]); ss

   user  system elapsed
  0.024   0.000   0.023

> head(ans2)

     x y           v
[1,] R h -0.8196355
[2,] R h  1.1127891
[3,] R h -0.2415036
[4,] R h  1.1443273
[5,] R h  1.4476976
[6,] R h  1.1084013

> dim(ans2)

[1] 14793      3

> identical(ans1$v,ans2$v)

[1] TRUE
```

1.2 Grouping

This is a repeat of the test in section 2 of the Introduction vignette. The syntax is explained there.

```
> ttt=system.time(ans1 <- tapply(DF$v,DF$x,sum)); ttt

   user  system elapsed
 17.725   1.052  18.853

> head(ans1)

        A         B         C         D         E         F
 670.5696 1339.5182 -456.3565 -954.6692  796.7028   33.4407

> sss=system.time(ans2 <- DT[,sum(v),by=x]); sss

   user  system elapsed
  0.464   0.180   0.645

> head(ans2)
```

```
      x        V1
[1,] A  670.5696
[2,] B 1339.5182
[3,] C -456.3565
[4,] D -954.6692
[5,] E  796.7028
[6,] F   33.4407

> identical(as.vector(ans1), ans2$V1)

[1] TRUE
```

1.3 Test 3

1.4 Test 4

1.5 Test 5

2 Summary table

```
> ans

        base data.table times faster
==     13.574      0.023          590
tapply 18.853      0.645           29

> toLatex(sessionInfo())
```

- R version 2.14.0 (2011-10-31), i686-pc-linux-gnu

- Locale: LC_CTYPE=en_GB.UTF-8, LC_NUMERIC=C, LC_TIME=en_GB.UTF-8, LC_COLLATE=C, LC_MONETARY=en_GB.UTF-8, LC_MESSAGES=en_GB.UTF-8, LC_PAPER=C, LC_NAME=C, LC_ADDRESS=C, LC_TELEPHONE=C, LC_MEASUREMENT=en_GB.UTF-8, LC_IDENTIFICATION=C

- Base packages: base, datasets, grDevices, graphics, methods, stats, utils

- Other packages: data.table~1.7.2

- Loaded via a namespace (and not attached): tools~2.14.0